

LATENT FIXTURE INFERENCE: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study tool use and manipulation. The motivating bet is: Infer unseen fixtures and environmental supports from action resistance patterns.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For tool use and manipulation, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026a;b;c;d;e; 2025). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.502	0.498
observed_only_wam	hidden_mode_shift	0.160	0.840
augmentation_as_data_wam	nominal	0.605	0.395
augmentation_as_data_wam	hidden_mode_shift	0.250	0.750
scalar_uncertainty_wam	nominal	0.710	0.290
scalar_uncertainty_wam	hidden_mode_shift	0.325	0.675
proposed_branch_mechanism	nominal	0.743	0.257
proposed_branch_mechanism	hidden_mode_shift	0.522	0.478

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Latent Fixture Inference argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Everydayvla: A vision-language-action model for affordable robotic manipulation, 2025. <http://arxiv.org/abs/2511.05397v1>.
- Grasp to act: Dexterous grasping for tool use in dynamic settings, 2026a. <http://arxiv.org/abs/2602.20466v1>.
- Pi-vla: Adaptive symmetry-aware decision-making for long-horizon vision-language-action manipulation, 2026b. <https://doi.org/10.3390/sym18030394>.
- A state-adaptive koopman control framework for real-time deformable tool manipulation in robotic environmental swabbing, 2026c. <https://doi.org/10.1002/adrr.202500142>.
- Plato: Planning with llms and affordances for tool manipulation, 2026d. <https://doi.org/10.1007/s10846-026-02392-y>.
- Langforce: Bayesian decomposition of vision language action models via latent action queries, 2026e. <http://arxiv.org/abs/2601.15197v7>.